

Zhengchun has 10+ years of experience conducting research and development around Data, Artificial Intelligence and High Performance Computing system. Currently, works as a Senior Machine Learning Scientist at AWS AI Labs. Specifically, he works at the intersection of machine learning and large scale distributed systems towards cloud-based Autonomous Data Warehouse.

## RESEARCH & ENGINEERING EXPERIENCE

- Amazon Web Service (AWS), AI Labs** Aug. 2022 — present  
*Senior Applied Scientist (Feb. 2023), Machine Learning Scientist II (Aug. 2022)*
  - Tech Lead of Machine Learning for AI-driven Autonomous Data Warehouse (Redshift Serverless)
  - Machine Learning in Database Engine for Workload Management, AI for Systems.
- Argonne National Laboratory, Data Science and Learning division** Aug. 2019 — Aug. 2024  
*on Sabbatical leave (Aug. 2022), Computer Scientist (Tenured, May 2022), Assistant Computer Scientist (Sep. 2019)*
  - Principal Investigator (PI) or co-PI of 10+ NSF/DOE research projects totaling more than \$12M.
  - Applied Artificial Intelligence for Science.
  - HPC/Systems for Artificial Intelligence for Science applications.
  - Data and Machine Learning for HPC System(characterize, predict and optimize).
- University of Chicago, Computation Institute** Mar. 2018 — Aug. 2019  
*Research Scientist*
  - Scalable cyber-infrastructure for smart/autonomous science ecosystems.
  - Performance modeling and characterization of High Performance Computing system.
- Argonne National Laboratory, Mathematics and Computer Science division** Sep. 2016 — Mar. 2018  
*Postdoctoral Researcher*
  - High Performance File Transfer over Wide-Area Science Network.
  - Modeling, Simulation and Optimization for scientific workflows over distributed infrastructures.
- Software Engineer (Part-time)** May 2010 — Jun. 2013  
*Xi'an FengLiTong Electronic, China*
  - Board Support Package development for  $\mu$ C/OS-II real-time operating system.
  - Developed a backend communication service for million-level vehicle-data-recorders (IoT).
- Founder** Sep. 2010 — Jul. 2013  
*Embedded System Technology Consulting & Outsourcing LLC, China*
  - 10+ cases, three prototypes went into mass production.

## SELECTED AWARDS, GRANTS & HONORS

- Best Paper award at XLOOP'21, FGCS'21, MLN'19, MLN'18, and TRIDENTCOM'18  $\times 5$  2021, 2019, 2018
- Impact Argonne award for notable achievement in Innovation  $\times 2$  May 2020, July 2022
- Top Winner of the Technology Challenge at the ACM/IEEE Supercomputing Conference (SC'19) Nov. 2019
- Pacesetter award by Argonne National Laboratory Jan. 2019

## EDUCATION

- **PhD., Computer Science**, Autonomous University of Barcelona, Spain Jul. 2016
  - With Honors(Cum Laude); Extraordinary Doctorate Award(1 out of 5 thesis).
- **MSc., Guidance, Navigation and Control**, Northwestern Polytechnical University, China Apr. 2013
- **BSc., Aircraft Manufacturing Engineering**, Northwestern Polytechnical University, China Jun. 2010
  - National Scholarship(award to top 0.2% of outstanding students) $\times 2$ , 2008 and 2009.

## ACTIVITIES

- Professional Membership: Association for Computing Machinery (ACM).
- Editorship: Journal of Future Generation Computer Systems (FGCS), 2020-2023, Impact Factor: 7.2.
- Workshop Co-Chair: SRMPDS 2017, 2018, 2019, 2020; AI-Science'19.
- Technical Program Committee: ICDS'19; DAAC'17-19; DLS'19; ICDCS'20; SC'20; HPCC'20; HiPC'22; SC'23; IPDPS'23,24.
- Journal Reviewer: MDPI Sensors; Elsevier FGCS, JOCS; IEEE Access, TPDS.

Papers on **Machine Learning and Systems**

1. Alexander van Renen, Dominik Horn, Pascal Pfeil, Kapil Eknath Vaidya, Wenjian Dong, Murali Narayanaswamy, **Zhengchun Liu**, Gaurav Saxena, Andreas Kipf, Tim Kraska. *Why TPC is not enough: An analysis of the Amazon Redshift fleet*. VLDB'24.
2. Vikram Nathan, Vikramank Singh, **Zhengchun Liu**, Mohammad Rahman, Andreas Kipf, Dominik Horn, Davide Pagano, Gaurav Saxena, Balakrishnan Narayanaswamy, Tim Kraska. *Intelligent Scaling in Amazon Redshift*. SIGMOD'24.
3. Ziniu Wu\*, Ryan Marcus, **Zhengchun Liu**, Parimarjan Negi, Vikram Nathan, Pascal Pfeil, Gaurav Saxena, Mohammad Rahman, Balakrishnan Narayanaswamy, Tim Kraska. *Stage: Query Execution Time Prediction in Amazon Redshift*. SIGMOD'24.
4. **Zhengchun Liu**, Rajkumar Kettimuthu, Michael E. Papka, Ian Foster. *FreeTrain: A Framework to Utilize Unused Supercomputer Nodes for Training Neural Networks*. IEEE/ACM International Symposium in Cluster, Cloud, and Grid Computing (CCGrid'23).
5. Ahsan Ali\*, Hemant Sharma, Rajkumar Kettimuthu, Peter Kenesei, Dennis Trujillo, Antonino Miceli, Ian Foster, Ryan Coffee, Jana Thayer, **Zhengchun Liu**. *fairDMS: Rapid Model Training by Data and Model Reuse*. 2022 IEEE International Conference on Cluster Computing. [arXiv:2204.09805](https://arxiv.org/abs/2204.09805)
6. Joaquin Chung, Wojciech Zacherek, AJ Wisniewski, **Zhengchun Liu**, Tekin Bicer, Rajkumar Kettimuthu and Ian Foster. *SciStream: Architecture and Toolkit for Data Streaming between Federated Science Instruments*. ACM HPDC'2022.
7. **Zhengchun Liu**, Ahsan Ali\*, Peter Kenesei, Antonino Miceli, Hemant Sharma, Nicholas Schwarz, Dennis Trujillo, Hyunseung Yoo, Ryan Coffee, Ryan Herbst, Jana Thayer, Chun Hong Yoon, Ian Foster. *Bridge Data Center AI Systems with Edge Computing for Actionable Information Retrieval*. XLOOP@SC'21. *Best Paper Awarded*.
8. **Zhengchun Liu**, Rajkumar Kettimuthu, Joaquin Chung, Rachana Ananthakrishnan, Michael Link, Ian Foster. *Design and Evaluation of a Simple Data Interface for Efficient Data Transfer Across Diverse Storage*. ACM Transactions on Modeling and Performance Evaluation of Computing Systems, 2021 Vol. 6, No.1.
9. Papadimitriou, George, Cong Wang, Karan Vahi, Rafael Ferreir da Silva, Anirban Mandal, **Zhengchun Liu**, Rajiv Mayania, Mats Rynge, Mariam Kiran, Vickie E. Lynch, Rajkumar Kettimuthu, Ewa Deelman, Jeffrey S. Vetter, Ian Foster. *End-to-End Online Performance Data Capture and Analysis for Scientific Workflows*. Future Generation Computer Systems, Vol. 117, 2021, Pages 387-400, *2021 Best Paper Award*.
10. Tirthak Patel, **Zhengchun Liu**, Rajkumar Kettimuthu, Paul Rich, Bill Allcock, Devesh Tiwari. *Job Characteristics on Large-Scale Systems: Long-Term Analysis, Quantification, and Implications*. The International Conference for High Performance Computing, Networking, Storage, and Analysis (SC'20).
11. **Zhengchun Liu**, Ryan Lewis\*, Rajkumar Kettimuthu, Kevin Harms, Philip Carns, Nageswara Rao, Ian Foster and Michael Papka. *Characterization and Identification of HPC Applications at Leadership Computing Facility*. International Conference on Supercomputing (ICS'20).
12. Yuanlai Liu\*, **Zhengchun Liu**, Rajkumar Kettimuthu, Nageswara Rao, Zizhong Chen and Ian Foster. *Data transfer between scientific facilities - bottleneck analysis, insights and optimizations*. IEEE/ACM International Symposium in Cluster, Cloud, and Grid Computing (CCGrid'19).
13. **Zhengchun Liu**, Rajkumar Kettimuthu, Prasanna Balaprakash, Nageswara S. V. Rao and Ian Foster. *Building a Wide-Area Data Transfer Performance Predictor: An Empirical Study*. International Conference on Machine Learning for Networking (MLN'18).
14. **Zhengchun Liu**, Rajkumar Kettimuthu, Ian Foster, Peter H. Beckman. *Towards a Smart Data Transfer Node*. Future Generation Computer Systems, 2018(89), Pages 10—18.
15. Rajkumar Kettimuthu, **Zhengchun Liu**, David Wheeler, Ian Foster, Katrin Heitmann, Franck Cappello. *Transferring a Petabyte in a Day*. Future Generation Computer Systems, 2018(88).
16. **Zhengchun Liu**, Rajkumar Kettimuthu, Ian Foster and Yuanlai Liu. *A comprehensive study of wide area data movement at a scientific computing facility*. IEEE 38th International Conference on Distributed Computing Systems (SNTA@ICDCS'18).
17. Rajkumar Kettimuthu, **Zhengchun Liu**, Ian Foster, Peter H. Beckman, Alex Sim, John Wu, Wei-keng Liao, Qiao Kang, Ankit Agrawal, and Alok Choudhary. 2018. *Toward Autonomic Science Infrastructure: Architecture, Limitations, and Open Issues*. The 1st Autonomous Infrastructure for Science workshop (AI-Science@HPDC'18).
18. **Zhengchun Liu**, Rajkumar Kettimuthu, Ian Foster and Nageswara S.V. Rao. *Cross-geography Scientific Data Transfer Trends and User Behavior Patterns*. International Symposium on High-Performance Parallel and Distributed Computing (HPDC'18).
19. **Zhengchun Liu**, Rajkumar Kettimuthu, Sven Leyffer, Prashant Palkar and Ian Foster. *A mathematical programming and simulation based framework to evaluate cyberinfrastructure design choices*. IEEE International Conference on eScience.
20. **Zhengchun Liu**, Prasanna Balaprakash, Rajkumar Kettimuthu and Ian Foster. *Explaining Wide Area Data Transfer Performance*. International Symposium on High-Performance Parallel and Distributed Computing (HPDC'17), 167-178.

## Papers on AI/HPC for Science

1. Anakha V Babu, Tao Zhou, Saugat Kandel, Tekin Bicer, **Zhengchun Liu**, William Judge, Daniel J Ching, Yi Jiang, Sinisa Veseli, Steven Henke, Ryan Chard, Yudong Yao, Ekaterina Sirazitdinova, Geetika Gupta, Martin V Holt, Ian T Foster, Antonino Miceli and Mathew J Cherukara. ***Deep Learning at the Edge Enables Real-time Streaming Ptychographic Imaging***. Nature Communications 14, 7059 (2023).
2. Lipeng Wan, Jieyang Chen, Xin Liang, Ana Gainaru, Qian Gong, Qing Liu, Ben Whitney, Joy Arulraj, **Zhengchun Liu**, Ian Foster, Scott Klasky. ***RAPIDS: Reconciling Availability, Accuracy, and Performance in Managing Geo-Distributed Scientific Data***. International Symposium on High-Performance Parallel and Distributed Computing (HPDC'23).
3. Petro Junior Milan, Hongqian Rong, Craig Michaud, Naoufal Layad, **Zhengchun Liu**, Ryan Coffee. ***Enabling real-time adaptation of machine learning models at x-ray Free Electron Laser facilities with high-speed training optimized computational hardware***. Frontiers in Physics, Volume 10 - 2022.
4. Mike Kraus, Naoufal Layad, **Zhengchun Liu**, Ryan Coffee. ***EdgeAI: Machine learning via direct attached accelerator for streaming data processing at high shot rate x-ray free-electron lasers***. Frontiers in Physics, Volume 10 - 2022.
5. **Zhengchun Liu**, Rajkumar Kettimuthu, Ian Foster. ***Masked Sinogram Model with Transformer for ill-Posed Computed Tomography Reconstruction: a Preliminary Study***. [arXiv:2209.01356](https://arxiv.org/abs/2209.01356)
6. **Zhengchun Liu**, Hemant Sharma, Jun-Sang Park, Peter Kenesei, Antonino Miceli, Jonathan Almer, Rajkumar Kettimuthu and Ian Foster. ***BraggNN: Fast X-ray Bragg Peak Analysis Using Deep Learning***. IJCrJ, Vol. 9, No. 1, 2022.
7. Aniket Tekawade, **Zhengchun Liu**, Peter Kenesei, Tekin Bicer, Francesco De Carlo, Rajkumar Kettimuthu, Ian Foster. ***3D Autoencoders For Feature Extraction In X-ray Tomography***. 2021 IEEE International Conference on Image Processing.
8. Jiali Wang, **Zhengchun Liu**, Ian Foster, Won Chang, Rajkumar Kettimuthu, Rao Kotamarthi. ***Fast and accurate learned multiresolution dynamical downscaling for precipitation***. Journal of Geoscientific Model Development.
9. Selin Aslan, **Zhengchun Liu**, Viktor Nikitin, Tekin Bicer, Sven Leyffer, Doga Gursoy. ***Joint Ptycho-Tomography with Deep Generative Priors***. Machine Learning Science and Technology, 2021, Vol. 2, No.4
10. Ziling Wu\*, Tekin Bicer, **Zhengchun Liu**, Vincent De Andrade, Yunhui Zhu, Ian T. Foster. ***Deep Learning-based Low-dose Tomography Reconstruction with Hybrid-dose Measurements***. AI4S@SC'20.
11. **Zhengchun Liu**, Tekin Bicer, Rajkumar Kettimuthu, Doga Gursoy, Francesco De Carlo and Ian Foster. ***TomoGAN: Low-Dose Synchrotron X-Ray Tomography with Generative Adversarial Networks***. Optical Society of America A, 2020, Vol. 37, No. 2.
12. Vibhatha Abeykoon\*, **Zhengchun Liu**, Tekin Bicer, Rajkumar Kettimuthu, Geoffrey Fox and Ian Foster. ***Scientific Image Restoration Anywhere***. XLOOP @SC'19.
13. **Zhengchun Liu**, Tekin Bicer, Rajkumar Kettimuthu and Ian Foster. ***Deep Learning Accelerated Light Source Experiments***. IEEE/ACM Deep Learning on Supercomputers DLS@SC'19.
14. **Zhengchun Liu**, Dolores Rexachs, Francisco Epelde, and Emilio Luque. ***A simulation and optimization based method for calibrating agent-based emergency department models under data scarcity***. Computers & Industrial Engineering, 2017.
15. Xueping Zhu, **Zhengchun Liu** and Jun Yang. ***Model of Collaborative UAV Swarm Toward Coordination and Control Mechanisms Study***. 2015 International Conference on Computational Science (ICCS'15), Vol 51, 493-502.
16. **Zhengchun Liu**, Eduardo Cabrera, Manel Taboada, Francisco Epelde, Dolores Rexachs and Emilio Luque. ***Quantitative Evaluation of Decision Effects in the Management of Emergency Department Problems***. International Conference on Computational Science (ICCS'15), Vol 51, Pages 433-442.
17. **Zhengchun Liu**, Eduardo Cabrera, Dolores Rexachs and Emilio Luque. ***A Generalized Agent-Based Model to Simulate Emergency Departments***. The 6th International Conference on Advances in System Simulation (SIMUL'14).

– Last updated on August 27, 2024